

HLT performance studies

in preparation for Run 3

Rafał Bielski

on behalf of the Trigger Core Software Team

Software Coordination meeting 29/10/2020



Outline

- ▶ Overview of HLT performance features and needs
 - ▶ Partially repeated from [SPOT meeting 14/09](#)
- ▶ HLT performance measurement and monitoring tools
- ▶ Recent results

HLT vs offline reconstruction

Similarities

- ▶ Both are part of Athena(MT) – a lot of common components between them
- ▶ Both use GaudiHive Avalanche Scheduler for Run 3

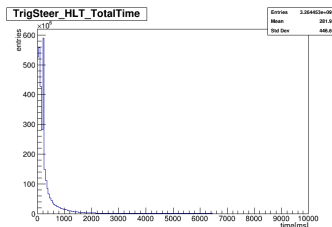
Differences

- ▶ Reco performance limited by memory needs, HLT by algorithm processing time and upstream I/O
- ▶ Run-3 HLT uses the same Scheduler but paired with complex HLT Control Flow
 - Partial event data processing (**event views**) and **early rejection** are the core of HLT design
- ▶ Very different and varied event-to-event CPU cost (see next slide)
- ▶ Main motivation for HLT athenaMT migration was **not memory saving**, but:
 - Closer integration with offline for easier maintenance and common features
 - Future-proofing – more modern code, potentially more flexible for future architectures
- ▶ HLT event processing involves Athena + [HLT MPPU](#) + [DCM](#) (online) / [DcmEmulator](#) (offline)
 - Possible to run HLT with `athena.py`, but `athenaHLT.py` (including HLT MPPU + DcmEmulator) better corresponds to online processing at P1

HLT performance needs

Purpose of the HLT: process events with very high input rate and select a small fraction of interesting ones

- ▶ 80–99% events are **rejected** (depending on streams and prescales) – these are never seen offline but constitute the majority of events in HLT and the performance needs to be optimised for them
- ▶ Most events take **<0.5 second** to process, some may take many seconds
- ▶ Performance needs constantly changing (also during a run) from:
 - ▷ Luminosity and pile-up
 - ▷ L1 rate and detector performance
 - ▷ Menu and prescales, streaming strategy



HLT event processing time online histogram, run [360026](#)

(the peak at 250-300 ms is from a "time burner" running in standby before collisions)

HLT performance needs

The most important HLT performance metric is **total event throughput**

- ▶ HLT farm needs to sustain L1 rate up to 100 kHz
- ▶ Throughput inversely proportional to average event processing time
- ▶ I/O generally fast at P1 (using only small fragments of full event), but need to watch the data request rates to avoid hitting ROS/network limitations
- ▶ Menu/prescale changes allow to increase throughput but at the cost of physics
- ▶ Higher throughput for a given menu = more space for lower-threshold / more complex selections = more interesting events for physics analysis
- ▶ No throughput gain from MT over MP by definition, but potentially lower memory needs allow to buy more CPUs*

Run-3 HLT is **MP+MT hybrid** opening a large parameter space for optimisation

- ▶ Need to find the N forks / N threads / N slots combination allowing the highest event throughput
- ▶ For two similarly-performing configurations, the one with lower memory usage would be preferred
- ▶ Need to take system stability and error-handling consequences into account

*Not necessarily possible with IT tender, will need to see next year

HLT performance measurement and monitoring

1) Trigger **cost monitoring**

- ▶ Collects execution time data for every algorithm in a fraction of events
- ▶ Allows to estimate CPU cost of each algorithm and predict CPU cost of the full menu with given prescales – essential for menu design and prescale choices
- ▶ Extensive offline analysis and visualisation tools (in progress for Run-3 framework)
- ▶ Results presented at atlas-trig-cost.cern.ch (example below)

Display Algorithm Summary

Home → Directory: Technical-and-M-weeks → Run: T13-PhysicsP1-1Forks-12Threads-12Slots 383763 → Range: LumiBlock 00000 → Summary: HLT Algorithm

Trigger Expert: [New Processing Request] [View P1 Logs] [View Request Logs] [View Installed SW]

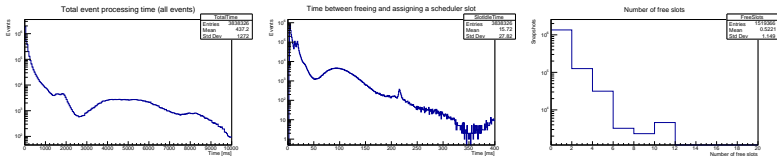
Table parsed from CSV file: [Technical-and-M-weeks/costMonitoring_T13-PhysicsP1-1Forks-12Threads-12Slots_383763/csv/Table_Algorithm_HLT_LumiBlock_00000.csv](#)

Algorithm Summary										
50 of 9 Page 1 of 9 Displaying 1 to 50 of 427 items										
Name	Raw Active Events	Active Events	Calls/Event	Calls > 1000 ms	Event Rate [Hz]	Call Rate [Hz]	Alg Total Time [s]	Alg Total Time [H]	Alg Total Time/Call [ms]	Alg Total Time/Event [ms]
TrigFastTrackFinder_jet	4294	4.294e+04	1	4.294e+04	51.12	51.12	2e+05	54.51	4663	4663
HLTCaloClusterMakerFS	5398	5.398e+04	1	5.398e+04	64.26	64.26	8.205e+04	22.36	1529	1529
PFAlgorithm_if	2962	2.962e+04	1	0	35.26	35.26	9527	2.597	323.8	323.8
TrigMuonCombinedAlg_Rol	1318	1.318e+04	1.179	2610	15.69	18.5	5427	1.479	346.6	408.7
HLTCaloCellMakerFS	5398	5.398e+04	1	100	64.26	64.26	5394	1.47	98.89	98.89
TrigCaloClusterMakerMT_topoLC	1364	1.364e+04	1.205	2910	16.24	19.56	3973	1.083	240.2	289.3
InDetSCT_Clusterization_jet	4294	4.294e+04	1	0	51.12	51.12	3901	1.063	90.81	90.81
TrigMuPatTrackBuilder_Rol	8158	8.158e+04	1.031	0	97.12	100.1	3518	0.9589	41.49	42.76
TrigMuonCandidateAlg_Rol	8158	8.158e+04	1.031	0	97.12	100.1	3092	0.8427	42.96	44.54
InDetPixelClusterization_jet	4294	4.294e+04	1	0	51.12	51.12	3067	0.8414	72.31	72.31
PFTrackSelector_if	2962	2.962e+04	1	0	35.26	35.26	2972	0.8102	100.2	100.2
InDetTrackerSpacePointFinder_jet	4294	4.294e+04	1	0	51.12	51.12	2918	0.7954	67.7	67.7
HLTCaloClusterCalibratorLCFS	4294	4.294e+04	1	0	51.12	51.12	2915	0.7946	67.8	67.8

HLT performance measurement and monitoring

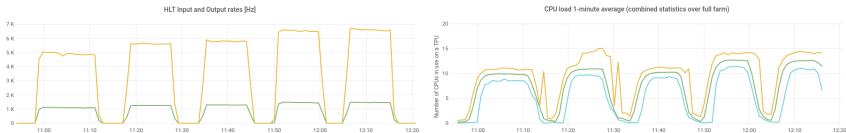
2) Online monitoring histograms

- ▶ A selection of histograms with simple metrics allow to evaluate overall HLT farm performance in real time
- ▶ Allow to narrow down performance limitations and other problems in real time
- ▶ Very useful also to analyse offline and compare configurations



3) Online monitoring data in PBEAST/Grafana

- ▶ Monitors huge amount of information at P1
- ▶ Most relevant for performance: rates and system load at different stages of TDAQ chain, mem/cpu



HLT performance measurement and monitoring

4) Dedicated offline studies (running athenaHLT* + perf tools)

- ▶ Collect prmon data, online histograms, cost data from different configurations to compare
 - ▶ Useful to debug specific issues or to test many configurations on small samples
 - ▶ Faster turnaround and easier to automatise than P1 running
 - ▶ athenaHLT scaling plots shown later are made using this approach
 - ▶ Close to online, but not identical – main differences in I/O (slower offline) and available resources (1 machine vs full HLT farm)
- ▶ Can also use Intel VTune to hunt for specific problems
 - ▶ Useful in hunting for fairly obvious and large issues
 - ▶ Limited usefulness otherwise as it cannot handle longer jobs with many threads/processes and software sampling is capped at 10 ms
 - ▶ Requires a lot of expertise and thought to interpret results case by case for anything other than "top 10 hot spots"

* athenaHLT.py emulates online running by reusing/emulating parts of the TDAQ software that steer Athena execution

Results introduction

The following slides present performance results from **athenaHLT offline** studies and from **P1 measurements** in Technical Runs

- ▶ We use labels "F-T-S" where F = N forks, T = N threads per fork, S = N slots per fork
 - ▶ N-1-1 is pure MP configuration, 1-N-M is pure MT, others are hybrid
- ▶ athenaHLT studies of wider phase-space are shown first
 - ▶ Chronologically, the P1 measurements were first and the offline ones were a follow-up
- ▶ Discovered a major flaw in data extraction procedure
 - ▶ Used **TH1::GetMean** from online histograms
 - ▶ Overflow is ignored by default and may introduce large error (shift)
 - ▶ Particularly for steeply falling distributions with long tails
 - ▶ Timing distributions and thus throughput measurement very affected
- ▶ Throughput information shown today extracted with alternative methods not using histograms
 - ▶ Detailed data separating event processing and I/O times available only from histograms
 - ▶ Plan to use extensible-axis histograms in the future
- ▶ We have lots of detailed in-job monitoring not shown today to help us understand the behaviour
 - ▶ For example data from SchedulerMonSvc presented in [SPOT meeting 14/09](#)

Scaling with athenaHLT

Release: Athena, 22.0.18

Trigger menu: LS2_v1 aka Dev_pp_run3_v1 without streamers (240 chains)

Data: 11.2k EnhancedBias stream events, run 360026 (late 2018), with original L1 prescales, decompressed

Machine: pc-tbed-pub-29, Xeon E5-2620 v4 2.10 GHz, 16 physical cores (2 sockets x 8 cores), 32 logical cores (2 threads per physical core), 32 GB RAM

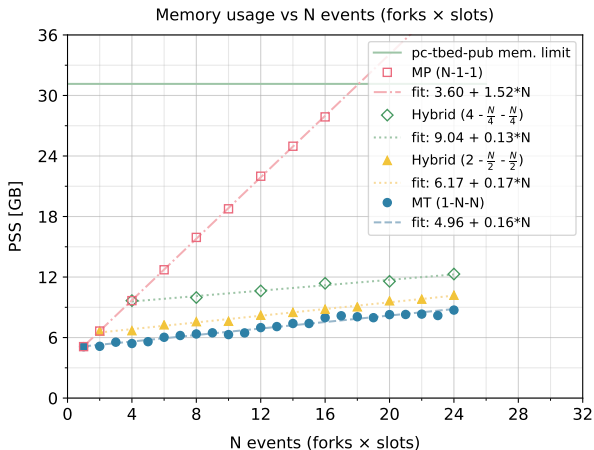
Configurations: threads=slots with 1, 2, 4 forks and pure MP, up to 24 total threads

Command:

```
prmon -i 3 -- athenaHLT.py \  
--nprocs=${nF} --threads=${nT} --concurrent-events=${nT} \  
-c "doStreamingSlice=False;doMonitorSlice=False;doBeamspotSlice=False;enableSchedulerMon=True;" \  
-R 360026 -f /scratch/rbielski/large.decompressed.data._0001.data \  
TriggerJob0pts/runHLT_standalone.py
```

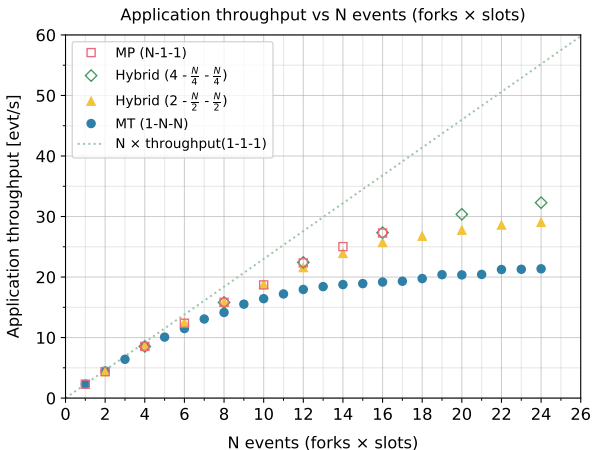
Disclaimer: we observe some MT errors in 22.0.18, mainly [ATR-22141](#), but these are rare and shouldn't affect the measurements

Scaling with athenaHLT – memory



- ▶ Expected behaviour: linear scaling for all configs with MT slope much smaller
- ▶ Need 5.1 GB for first fork, then 100–200 MB per slot and 1.5 GB per fork
- ▶ At 24 events factor 4.5 difference between pure MP and pure MT


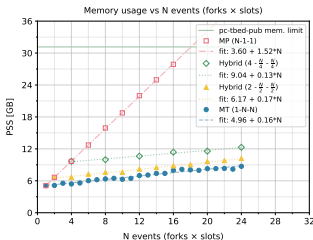
Scaling with athenaHLT – throughput



- ▶ Fairly good scaling until 8 events, then processing starts to slow down considerably
- ▶ Pure MT configuration much slower than pure MP at high thread/process counts
- ▶ Both trends partially due to hardware / low-level effects, but software effects are under investigation



Comparison to offline reconstruction

HLT performance as measured with athenaHLT on pc-tbed-pub:




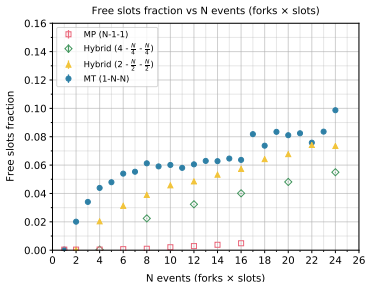
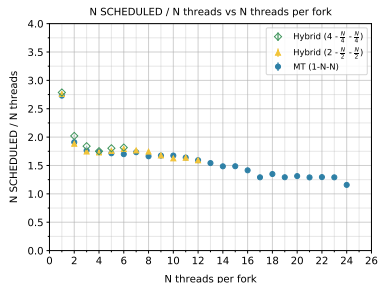
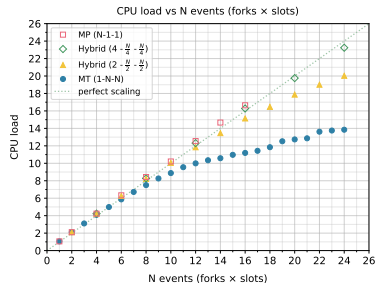
Offline reco performance on data as shown [last week](#) by Johannes:

RELEASE 22 RECO_TF DATA18 MEMORY AND EVENT THROUGHPUT SCALING



Very tempting to do this comparison, though the jobs are very different and probably mechanisms of saturation are also different to some extent

Teaser: other interesting plots



P1 performance results

Release: Athena, 22.0.18

Trigger menu: "Dev" (240 chains) and "PhysicsP1" (54 chains)

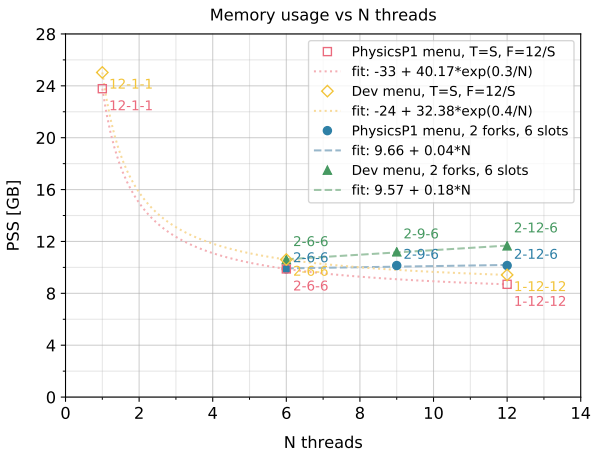
Data: 2.5k EnhancedBias stream events, run 360026 (late 2018), with original L1 prescales, preloaded into ROS at P1 and looped over many times ($\sim 3\text{M}$ events per run)

Machines: 5 racks, 39 TPUs each ($\sim 10\%$ of the farm), each TPU with Xeon E5-2660 v4 2.0 GHz, 28 physical cores (2 sockets \times 14 cores), 56 logical cores (2 threads per physical core), 65 GB RAM

Configurations: Tested **5 ways of processing 12 events in parallel** on each node (F-T-S = forks-threads-slots) – here in order from the least to the most MT:

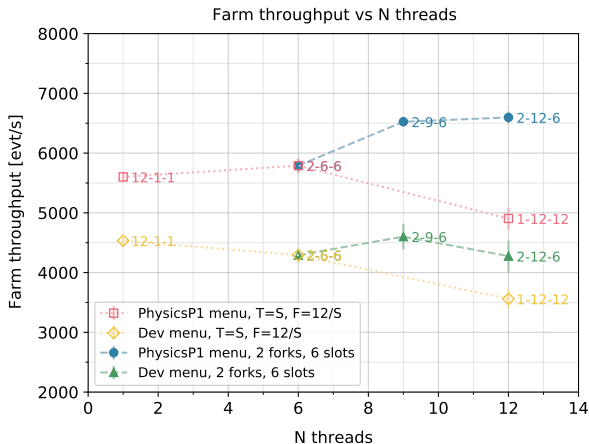
- ▶ 12-1-1 - pure MP, no MT, like in Run 2
- ▶ 2-6-6 - hybrid with threads = slots (12 threads in total)
- ▶ 2-9-6 - hybrid with threads = $1.5 \times$ slots (18 threads in total)
- ▶ 2-12-6 - hybrid with threads = $2 \times$ slots (24 threads in total)
- ▶ 1-12-12 - pure MT, no MP, threads = slots

Scaling at P1 – memory



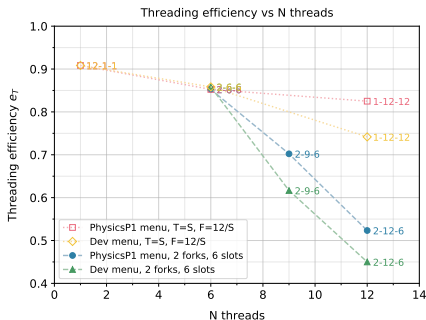
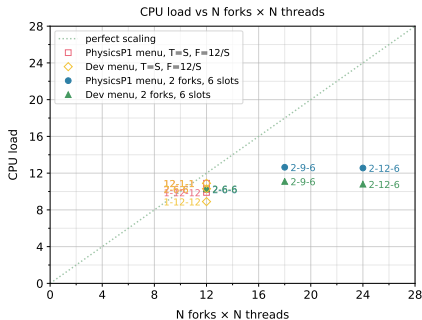
- ▶ Expected: MT uses much less memory than MP
- ▶ Expected: Dev menu (more chains) needs slightly more memory than PhysicsP1
- ▶ Interesting but not a problem: Adding threads per event requires a little bit more memory?

Scaling at P1 – throughput



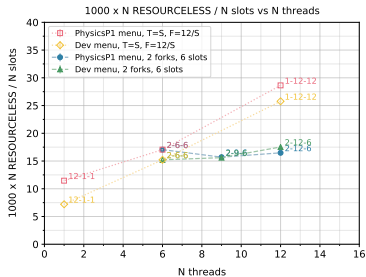
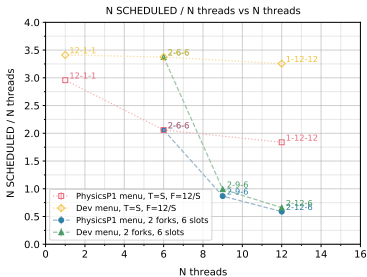
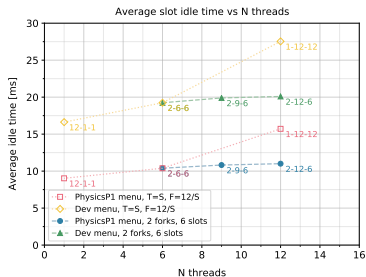
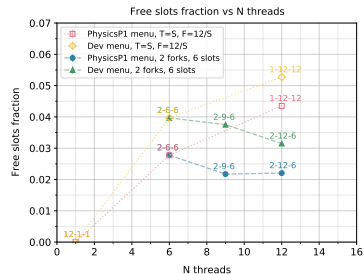
- ▶ At threads=slots, MT configurations are slower than MP (also seen offline)
- ▶ Adding a few more threads (1.5 per event on average) speeds up the processing
- ▶ The speed-up saturates quickly, no improvement with 2 threads per event

Scaling at P1 – CPU load



- ▶ Extra threads beyond 1.5 thread per event remain idle – no gain from adding more
- ▶ Mechanism is understood – mostly sequential Control Flow required for early rejection combined with a few long critical algorithms (see also the [SPOT meeting slides](#))
- ▶ The optimal threads-to-slots ratio depends on menu, but is always slightly above 1
- ▶ We can monitor scheduler under-allocation in real time online thanks to SchedulerMonSvc
- ▶ The inefficiency of MT at threads=slots with respect to MP is under investigation ([ATR-22112](#))

Teaser: other interesting plots



Summary

- ▶ Presented HLT performance needs, monitoring tools and results of P1 and offline scaling studies
- ▶ Extensive inventory of tools allows us to extract detailed information both in real time and post-run
- ▶ Memory scaling as expected, a lot can be saved in MT
 - ▷ Not so crucial for HLT, but useful input for Run-3 hardware purchase
- ▶ Pure MT at threads=slots is slower than pure MP
 - ▷ To some extent natural due to hardware/low-level effects
 - ▷ Partially caused by I/O and in-event effects (locks, unclonable algs)
 - ▷ Investigating whether there is still space for software improvements
- ▶ Some throughput improvement possible with larger thread pool
 - ▷ The optimal configuration for HLT somewhere between 1 and 1.5 threads per event
- ▶ To define the best Run-3 configuration, need to take into account throughput, memory usage, error handling, stability
 - ▷ Most likely will run with hybrid mode with a few slots per fork and 1.2–1.5 threads per slot
- ▶ Timeline:
 - ▷ Implement all chains by the end of 2020
 - ▷ Validation / optimisation / CPU cost estimates in winter/spring 2021
 - ▷ Run-3 farm hardware choice and order in summer 2021

Past events:

- ▶ [May 2020 HLT Hackathon](#)
- ▶ [19/06/2020 Trigger Core Software meeting](#)
- ▶ Technical Run 10 [Jira ticket](#) and [performance plots](#)
- ▶ Technical Run 12 [Jira ticket](#) and [performance plots](#)
- ▶ [14/09/2020 SPOT meeting](#)
- ▶ [September 2020 HLT Hackathon](#)
- ▶ Technical Run 13 [Jira ticket](#) and [09/10/2020 Trigger Core Software meeting](#)
- ▶ [23/10/2020 Trigger Core Software meeting](#)

Future events:

- ▶ [Trigger Workshop](#), 9–13/11/2020
- ▶ [Technical Run 14](#), 16–20/11/2020